

Intorduction to Dynamic Programming

About me

- Carlos Gonzalez Oliver
- **carlos@ozeki.io** (<mailto:carlos@ozeki.io>)
- **carlosoliver.co** (<http://carlosoliver.co>)

My interests:

- Structural Bioinformatics
- Graph algorithms
- Representation learning

Some questions I ask:

- *How can we efficiently search through graph databases?*
- *What kinds of patterns can we discover in datasets of graphs?*
- *Can we model biological systems using efficient data structures to learn about how they function?*

Dynamic Programming (DP)



- Roger Bellman, 1950s, working at RAND institute.
- Picked the name "Dynamic Programming" to please his boss.

- Method for solving problems by constructing the solution bottom-up.

Applications

- Bioinformatics: protein and RNA folding, sequence alignment
- Speech recognition: Viterbi's algorithm
- Time series modeling: Dynamic time warping
- Search: string matching
- Scheduling: weighted interval
- Music: Beat tracking
- Routing: shortest path problems

Objectives

After this talk we should be able to:

- Recognize the **ingredients** needed to solve a problem with DP.
- Write down and execute some simple DP algorithms.
- Be able to implement a DP algorithm for a real-world problem (homework).

Basic Intuition

- What is:
 - $(1 + 1 + 1 + 1 + 1 + 1 = ?)$
- Now what is:
 - $(1 + 1 + 1 + 1 + 1 + 1 + 1 = ?)$

First ingredient: Optimal Substructures

The optimal structure can be built by combining optimal solutions to smaller problems.

E.g. Fibonacci numbers

- $(\text{Fib}(0) = 1)$
- $(\text{Fib}(1) = 1)$
- $(\text{Fib}(n) = \text{Fib}(n-1) + \text{Fib}(n-2))$

Second ingredient: Overlapping Substructures

The same subproblem's solution is used multiple times.

```

1  def fib(n):
2      if n == 0 or n == 1:
3          return 1
4      else:
5          return fib(n-1) + fib(n-2)

```

Call tree `fib(6)` :

Recursive solution is $\mathcal{O}(2^n)$.

An example *without* overlapping subproblems.

Binary Search

```

def binary_search(A, left, right, query):
    if low > high:
        return None

    mid = (left + right) // 2
    if A[mid] == query:
        return i
    elif A[mid] > query:
        return binary_search(A, left, mid-1, query)
    else:
        return binary_search(A, mid+1, right, query)

```

Call tree `binary_search(A=[15, 22, 32, 36, 41, 63, 75], left=0, right=4, query=75)` :

Case study: Weighted Interval Scheduling

- **Input:** (n) requests, labeled $\{1, \dots, n\}$. Each request has a start time (s_i) and an end time (f_i) , and a **weight** (v_i) .
- **Output:** a *compatible* subset (S) of $\{1, \dots, n\}$ that maximizes $(\sum_{i \in S} v_i)$.

\mathcal{S} is compatible iff all pairs of intervals in \mathcal{S} are non-overlapping.

- Uses: resource allocation for computer systems, optimizing course selection.

Example:

False start: greedy approach

- From "previous" lectures we know the greedy approach to the unweighted IS problem gives the optimal solution (i.e. pick item with earliest ending time)


Counterexample:

The weights force us to consider all possible subproblems (i.e. local choice is not enough).

A helper function

Let us sort all intervals by increasing *finish* time and let $\text{dp}(j)$ return the latest interval i that is still compatible with j .

Index



First ingredient: Optimal Substructure?

- Consider \mathcal{O}_n to be the *optimal* set of requests over all items \mathcal{I}_n .
- There are two cases for the last request, \mathcal{I}_n :

Case 1

Case 1

Implication?

Optimal Substructure

Let $\text{OPT}(j)$ be the total weight of the optimal over intervals up to j .

We can now write an expression for computing $\text{OPT}(j)$:

Ingredients

- ☒ Optimal Substructure
- ☐ Overlapping Subproblems

First algorithm


Now we can write down a recursive algorithm that gives us the maximum weight over $(1, \dots, n)$.

`compute_OPT(j) :`

Ingredient 2: Overlapping Substructures

Let's build the execution tree for our recursive algorithm on this example:

Index



Runtime:

Ingredients

- ✓ Optimal Substructure
- ✓ Overlapping Subproblems

Memoization

- Key idea in DP: remember solutions to sub-problems and iterate such that you already have the solution you need.

New algo: $M_compute_OPT(j)$

Runtime:

Are we done?


- Recall $OPT(j)$ is just a number, we want the set of intervals with the score $OPT(j)$... i.e. \mathcal{O}_j .
- **Traceback** is a key idea in DP. We reconstruct the solution backwards from the M array using the recurrence.
- **Obervation:** We know that an interval j belongs to \mathcal{O}_j if:

Now we just apply this rule going backwards from $j = n$ and collect optimal intervals until we reach 0.

Full Example

Index

1	$v_1 = 2$	$p(1) = 0$
2	$v_2 = 4$	$p(2) = 0$
3	$v_3 = 4$	$p(3) = 1$
4	$v_4 = 7$	$p(4) = 0$
5	$v_5 = 2$	$p(5) = 3$
6	$v_6 = 1$	$p(6) = 3$




Fill M and get \mathcal{O}_n :

Recap

- Dynamic Programming works well when we have a problem structure such that:
 - Combining sub-problems solves the whole problem
 - Sub-problems overlap
- We can often reduce runtime complexity from exponential to polynomial or linear.

General interest: RNA folding

- RNA molecules are essential to all living organisms.
- Knowing the sequence is easy but not so informative, knowing the structure is hard but tells us a lot about the molecule's function.





A Sketch of Nussinov's Algorithm

- First attempt at solving this problem



Ruth Nussinov, designed the first RNA folding algo in 1977.

- In CS terms, an RNA is a string on a 4-letter alphabet.
- An RNA **structure** is a pair of indices over the string with certain constraints.



RNA folding rules

We use some fairly realistic constraints on admissible structures:

1. Only A-U, U-A and C-G, G-U form pairs
2. Pairs shall not cross (nestedness).
3. Start and end of a pair should be separated by at least θ spaces (remove steric clashes).
4. The best structure is the one that forms the most pairs (stability). e.g each pair adds 1 to the score.

This lets us identify the problem structure needed for a DP solution:

Consider an optimal set of pairs \mathcal{O}_{ij} between two indices (i, j) , and the score of the solution $\text{OPT}(i, j)$. We have two cases for index j :

Case 1. j is not in \mathcal{O}_{ij} :

Note we introduce a new variable \rightarrow 2 dimensional DP.

Case 2. j is in \mathcal{O}_{ij} :

Ingredients?

- ☐ Optimal substructures

- Overlapping subproblems

From this we can build our recurrence that fills the table up to $(\mathrm{OPT}(1, L))$.

Bonus Questions: What is the runtime for filling the table?

If you want an exercise sheet to learn how to implement this send me an email carlos@ozeki.io .

